

[COMPANY LOGO]
Insert logo here

A NEW SOUTH TECHNOLOGIES PUBLICATION | JUNE 2026

THE AI GOVERNANCE PLAYBOOK

Deploying Responsible AI Without Compromising Audit Readiness

A practical guide for federal agencies navigating OMB M-26-04, NIST AI RMF, and responsible AI procurement.

NIST AI RMF
Compliant Framework

OMB M-26-04
Ready Deployment

FedRAMP
Authorization Ready

01
of 7

Introduction — Why AI Governance Cannot Be an Afterthought

Why AI Governance Cannot Be an Afterthought

Artificial Intelligence offers unprecedented opportunities for federal agencies — but it also introduces unique risks. Bias, opacity, data privacy violations, and mission-critical failures can trigger OIG investigations and public trust crises.

This Playbook provides a practical framework for deploying AI systems that are not only effective but also compliant with federal mandates.

Regulatory Mandates Covered

- NIST AI Risk Management Framework (AI RMF)
- OMB Memorandum M-26-04 (Unbiased AI Principles)
- Executive Order 14110 (Safe, Secure, Trustworthy AI)
- FedRAMP Authorization Requirements

By the end of this Playbook, you will have:

- ✓ A checklist for AI system assessment before deployment
- ✓ Templates for documenting model transparency and explainability
- ✓ Integration points for existing NIST RMF compliance workflows
- ✓ Strategies for continuous monitoring and risk management

02 of 7

Regulatory Landscape — What You Must Know Before Procurement

What You Must Know Before Procurement

Before your agency initiates any AI procurement or development program, these six regulatory frameworks define your compliance obligations. Non-compliance is not an option — it is an audit finding.

REGULATION	KEY REQUIREMENT	DEADLINE
NIST AI RMF	All AI systems handling federal data must follow Govern, Map, Measure, Manage functions	Ongoing
OMB M-26-04	Empirical evidence of bias mitigation prior to contract award; explainability for high-stakes decisions	Immediate
EO 14110	Safety testing standards; watermarks for AI-generated content; internal governance policies required	Ongoing
FedRAMP	Cloud-based AI services require full ATO authorization before deployment	Per Deployment
Privacy Act §552a	PII handling requires explicit consent and documented data retention schedules	Ongoing
CUI Protection	Training data containing classified info requires additional safeguards and access controls	If Applicable

“The Bottom Line: If you cannot explain *why* your AI made a decision, you cannot authorize it.

03 of 7

Pre-Deployment Checklist — AI System Assessment

AI System Assessment Checklist

Use this checklist before beginning any AI procurement or development project. Every unchecked item is a potential audit finding.

☐ DATA GOVERNANCE

- ☐ Training data provenance documented and vetted
- ☐ Data classification verified (CUI, PII, FOUO, Unclassified)
- ☐ Bias assessment completed across protected classes
- ☐ Consent mechanisms in place for personal data usage
- ☐ Data retention and deletion schedules defined

☐ MODEL TRANSPARENCY

- ☐ Model architecture documented for audit review
- ☐ Explainability layer implemented (SHAP, LIME, or equivalent)
- ☐ Decision logic traceable for individual outputs
- ☐ Known limitations documented in SSP
- ☐ Confidence scores displayed for all predictions

☐ SECURITY CONTROLS

- ☐ NIST 800-53 controls mapped to AI-specific risks
- ☐ Encryption at rest and in transit for all model artifacts
- ☐ Access controls defined for model training and inference
- ☐ Vulnerability scanning integrated into CI/CD pipeline
- ☐ Incident response plan includes AI-specific scenarios

☐ OPERATIONAL READINESS

- ☐ Human-in-the-loop process defined for high-risk decisions
- ☐ Monitoring dashboard for model drift and performance degradation
- ☐ Feedback mechanism for users to report errors or concerns
- ☐ Rollback procedures tested and documented
- ☐ Staff trained on AI system operation and ethics

Scoring Guide

- ✓ **All boxes checked** → Ready for Development
- ⚠ **Any box unchecked** → Pause and address the gap first

Black-Box Models Won't Pass Audit

Not all AI needs the same level of explainability. Classify your system using the matrix below to determine what documentation depth is legally required before authorization.

IMPACT LEVEL	EXAMPLE SYSTEMS	REQUIRED OUTPUT
HIGH	Benefits denial, parole, hiring decisions	Full feature-level explanation + human review before action taken
MEDIUM	Resource allocation, fraud detection	Summary explanation + confidence score
LOW	Email routing, document tagging	Category label only

Three Levels of Explainability Required

Local Explainability (Per Decision)

For each output, identify which input features contributed most. Must be understandable by a non-technical stakeholder.

Example: "Claim denied because income exceeds \$50,000 AND dependents < 2"

Global Explainability (Model-Wide)

Aggregate view of what factors drive the model overall. Used for bias detection and policy validation.

Example: "Income contributes 40% to decisions; age contributes 5%"

Temporal Explainability (Time-Based)

Tracking how explanations change as the model updates. Required for continuously learning models to detect drift in decision logic.

05
of 7

Vendor Evaluation — Don't Trust Claims, Verify Them

Don't Trust Vendor Claims — Verify Them

Every AI vendor will claim compliance, security, and fairness. Use these questions — not their marketing materials — as your standard of proof during procurement evaluation.

TECHNICAL CAPABILITIES

1. What algorithm was used to train the model? Why was this choice made?
2. Can you demonstrate how your system explains individual decisions?

3. How does the model handle edge cases or out-of-distribution inputs?
4. What is your approach to detecting and correcting bias during training?
5. Can we deploy a copy of the model within our network for validation?

COMPLIANCE SUPPORT

1. Will you help us map your controls to NIST 800-53 and AI RMF?
2. What audit logs do you provide? Are they immutable and tamper-proof?
3. Do you have FedRAMP authorization for your AI platform?
4. How often do you perform third-party security assessments?
5. What happens if an inspector finds a control violation in your system?

⚠️ RED FLAGS TO WATCH FOR

- ⚠️ Vague answers about data sourcing ("proprietary methods")
- ⚠️ Refusal to allow independent security audits
- ⚠️ No documentation for how model decisions were reached
- ⚠️ Claims of "bias-free" without empirical testing data

06
of 7

Implementation Roadmap — 12-Month Timeline

12-Month Implementation Timeline

Structured into six progressive phases, this roadmap moves your agency from initial inventory to full AI Authorization to Operate (ATO) — with governance embedded at every step.

MONTH	PHASE	KEY ACTIVITIES
1–2	Foundation	Appoint AI Governance Lead; inventory all existing AI/ML systems currently in use
3–4	Policy Development	Draft agency AI governance policy; define High-Stakes classification criteria
5–6	Pilot Program	Select one low-risk AI use case for a controlled pilot; implement monitoring dashboard
7–8	Scale Controls	Apply successful pilot controls to additional use cases; integrate into DevSecOps
9–10	3rd-Party Validation	Engage 3PAO for independent AI system assessment; address all findings
11–12	Full Authorization	Submit complete AI authorization package; achieve formal ATO with AI-specific conditions

Success Metrics

100%

AI Systems Mapped

All systems aligned to NIST AI RMF functions

Zero

Unexplained Decisions

No high-stakes outputs without traceable logic

<48h

Remediation Window

From anomaly detection to corrective action



Executive Dashboard

Monthly compliance posture visibility for leadership

07
of 7

Next Steps — Ready to Deploy AI Responsibly?

Ready to Deploy AI Responsibly?

Modernization without accountability isn't progress — it's risk. With a governed approach, you can deploy AI safely, transparently, and sustainably, and maintain the public trust your mission depends on.

Get a Free AI Governance Gap Analysis

Discover how your current AI initiatives measure up against OMB M-26-04 and NIST AI RMF requirements.

 **Email:** info@newsouthtechnologies.com

 **Web:** www.newsouthtechnologies.com

 **Phone:** [Insert Phone Number]

“We build the governance that makes AI innovation possible.”

— New South Technologies

© 2026 New South Technologies. All Rights Reserved.

This Playbook is provided for educational purposes only. Consult legal counsel for interpretation of regulatory requirements specific to your agency.